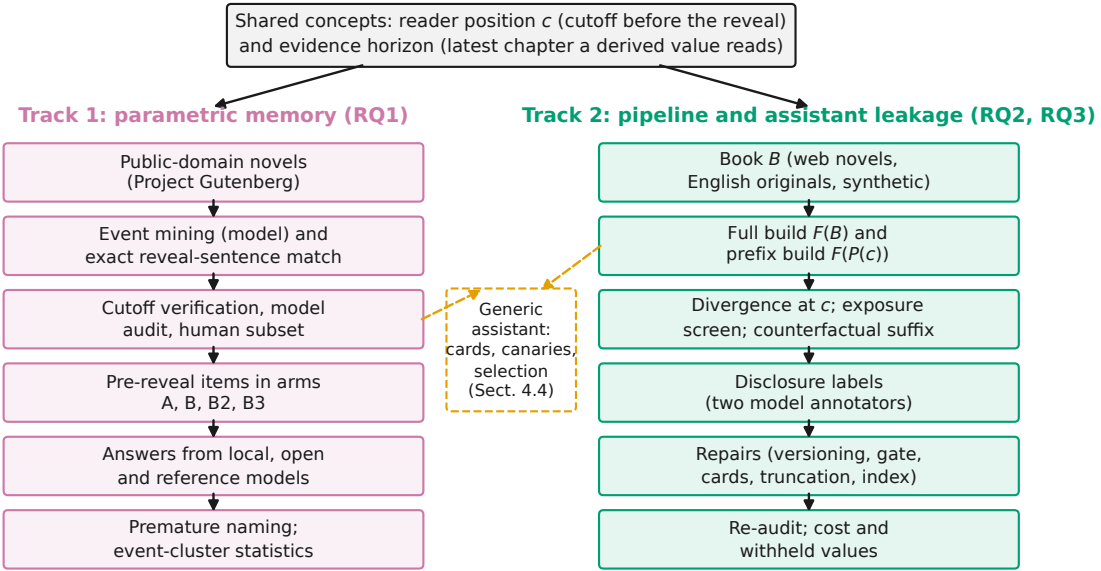
Shared concepts: reader position c (cutoff before the reveal) and evidence horizon (latest chapter a derived value reads)
Track 1: parametric memory (RQ1)
Track 2: pipeline and assistant leakage (RQ2, RQ3)
Public-domain novels (Project Gutenberg)
Event mining (model) and exact reveal-sentence match
Cutoff verification, model audit, human subset
Pre-reveal items in arms A, B, B2, B3
Answers from local, open and reference models
Premature naming; event-cluster statistics
Generic assistant: cards, canaries, selection (Sect. 4.4)
Book B (web novels, English originals, synthetic)
Full build $F(B)$ and prefix build $F(P(c))$
Divergence at c; exposure screen; counterfactual suffix
Disclosure labels (two model annotators)
Repairs (versioning, gate, cards, truncation, index)
Re-audit; cost and withheld values